

OS Homework-3 Salman Zafar

Ikhlas Ahmed Khan

27096

1. Why do we need push_up and pop_up on top of intr_on and intr_off?

Ans) We need push_up in acquire to disable all interrupts before entering the while loop inside acquire, which handles the locking. Similarly, we need pop_up after releasing the lock in the release function to re-enable interrupts. This is done to prevent any timer interrupt from occurring between the lock acquisition and its release. If a timer interrupt were to occur, it would itself acquire &tickslock to increment the ticks counter. If another user process acquires a lock during this stage, it would cause a deadlock. To prevent this, push_up and pop_up are used to disable and enable interrupts, respectively.

2. What is the difference between locks and semaphores usage-wise?

Ans) Locks are mainly useful in scenarios where the expected time a thread will spend waiting for the lock to be released is short. For example, if a process acquires a lock, it continuously checks in a while loop to see if the lock becomes available. For short critical sections where a lock is held briefly, this approach is efficient. However, if a process has to wait for a longer period, this busy waiting can be inefficient and wasteful. In contrast, semaphores are designed to avoid busy waiting in situations where the expected waiting time may be prolonged. In such cases, xv6 uses sleep/wakeup calls with semaphores, putting the waiting process into a sleep state. This allows the respective process to pause without consuming CPU resources, improving efficiency and resource management.

LOCK

locks can be used only for mutual exclusion.

A lock is a low-level synchronization mechanism.

SEMAPHORE

Semaphores can be used either for mutual exclusion or as a counting semaphore.

A semaphore is a signaling mechanism.

locks allow only one process at any given time to access the critical section.

locks can be wasteful if they are hold for a long time duration.

Only one thread is allowed at a time to acquire the lock and proceed with a critical section.

locks are very efficient because they are blocked only for a short period of time.

In locks, a process waiting for lock will keep the processor busy by continuously polling the lock.

locks are valid for only one process.

In locks, a process waiting for lock will instantly get access to a critical region as the process will poll continuously for the lock.

It is a busy wait process.

locks can have only two values – LOCKED and UNLOCKED

Semaphores allow more than one process at any given time to access the critical section.

In semaphore there is no resource wastage of process time and resources.

One or several threads are allowed to access the critical section.

Semaphores are held for a longer period of time. To access its control structure it uses spin lock.

In semaphore, a process is waiting for a semaphore to go into sleep to be woken up at any time and then try for the lock again.

Semaphores can be used to synchronize between different processes.

In semaphore, a process waiting for a lock might not get into the critical region as soon as the lock is free because the process would have gone to sleep and when it is woken up it will enter into the critical section.

It is a sleep wakeup process.

In semaphore, mutex will have value 1 or 0, but if used as counting semaphore it can have different values.

In uniprocessor system locks are not very useful because they will keep the processor busy every time while polling for the lock , thus disabling any other process from running.

In locks it is recommended to disable the interrupts while holding the locks.

Thread cannot sleep while waiting for the lock when it fails to get the lock, but it continues a loop of trying to get locked.

In a uniprocessor system semaphores are convenient because they don't keep the processor busy while waiting for the lock.

Semaphore can be locked with interrupt enabled.

Thread goes to sleep waiting for the lock when it fails to get the lock.

3. How many times in the code so far locks and semaphores have occurred?

Ans) In xv6, locks are acquired and released in various key areas, giving occurrences in 14 different files. These can be found in the following files: (note in brackets i have mentioned that how many times locks have occurred in the following codes)

1. sleeplock.c (3 times)

```
kernel/sleeplock.c

21 void
22 acquiresleep(struct sleeplock *lk)
23 {
24     acquire(&lk->lk);
30     release(&lk->lk);

33 void
34 releasesleep(struct sleeplock *lk)
35 {
36     acquire(&lk->lk);
40     release(&lk->lk);
48     acquire(&lk->lk);
50     release(&lk->lk);
```

2. bio.c(7)

```
kernel/bio.c

63     acquire(&bcache.lock);

68     b->refcnt++;
69     release(&bcache.lock);
70     acquiresleep(&b->lock);

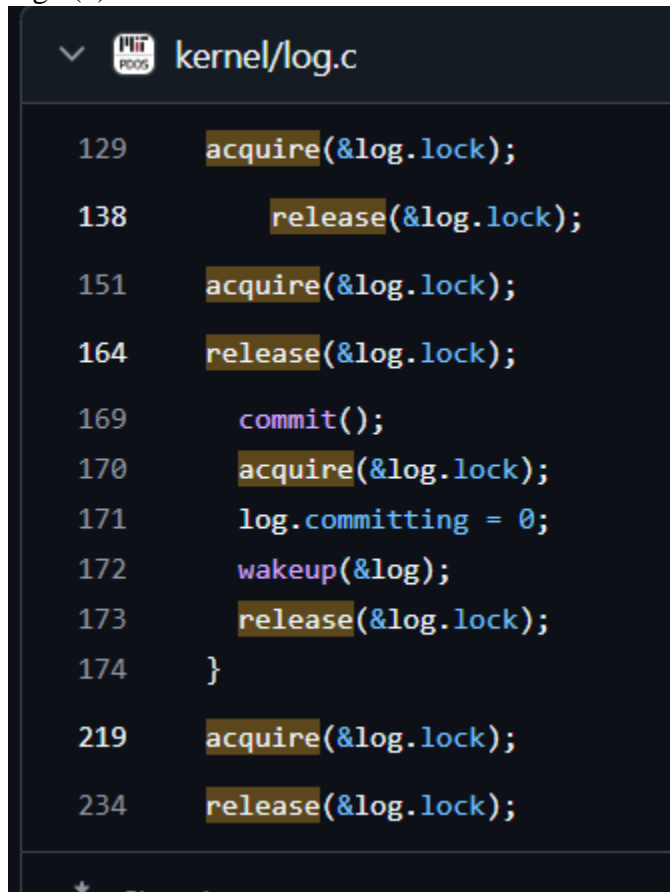
83     release(&bcache.lock);
84     acquiresleep(&b->lock);

114 // Release a locked buffer.
122     releasesleep(&b->lock);
123
124     acquire(&bcache.lock);
125     b->refcnt--;

136     release(&bcache.lock);

141     acquire(&bcache.lock);
143     release(&bcache.lock);
148     acquire(&bcache.lock);
150     release(&bcache.lock);
```

3. log.c(4)



```
kernel/log.c

129     acquire(&log.lock);
138         release(&log.lock);
151     acquire(&log.lock);
164     release(&log.lock);
169         commit();
170     acquire(&log.lock);
171     log.committing = 0;
172     wakeup(&log);
173     release(&log.lock);
174 }

219     acquire(&log.lock);
234     release(&log.lock);
```

4. pipe.c(5)

```
kernel/pipe.c

61     acquire(&pi->lock);
70     release(&pi->lock);
73     release(&pi->lock);
81
82     acquire(&pi->lock);
83     while(i < n){
84         if(pi->readopen == 0 || killed(pr)){
85             release(&pi->lock);
86             return -1;
100    release(&pi->lock);
112    acquire(&pi->lock);
115        release(&pi->lock);
128    release(&pi->lock);
```

5. file.c(4)



```
kernel/file.c

34  acquire(&ftable.lock);
38  release(&ftable.lock);
42  release(&ftable.lock);
49  {
50  acquire(&ftable.lock);
51  if(f->ref < 1)
53  f->ref++;
54  release(&ftable.lock);
55  return f;
64  acquire(&ftable.lock);
68  release(&ftable.lock);
74  release(&ftable.lock);
```


6. proc.c (16 locks)

```
26  // must be acquired before any p->lock.
97  acquire(&pid_lock);
100 release(&pid_lock);
115  acquire(&p->lock);
119  release(&p->lock);
131  release(&p->lock);
139  release(&p->lock);
254 release(&p->lock);
294  release(&np->lock);
315  release(&np->lock);
317  acquire(&wait_lock);
319  release(&wait_lock);
321  acquire(&np->lock);
323  release(&np->lock);
368  acquire(&wait_lock);
376  acquire(&p->lock);
381  release(&wait_lock);
397  acquire(&wait_lock);
405  acquire(&pp->lock);
413  release(&pp->lock);
556  // (wakeup locks p->lock),
557  // so it's okay to release lk.
558
558
559  acquire(&p->lock); //DOC: sleeplock1
560  release(lk);
```

7. kalloc.c (2 locks)

```
kernel/kalloc.c

58
59     acquire(&kmem.lock);
60     r->next = kmem.freelist;
61     kmem.freelist = r;
62     release(&kmem.lock);
63 }

73     acquire(&kmem.lock);
77     release(&kmem.lock);
```

8. sysproc.c (3)

```
kernel/sysproc.c

60     acquire(&tickslock);
64         release(&tickslock);
69     release(&tickslock);
86 {
87     uint xticks;
88
89     acquire(&tickslock);
90     xticks = ticks;
91     release(&tickslock);
92     return xticks;
```

9. console.c(3)

```
kernel/console.c

86     target = n;
87     acquire(&cons.lock);
88     while(n > 0){

92         if(killed(myproc())){
93             release(&cons.lock);
94             return -1;

124     release(&cons.lock);
138     acquire(&cons.lock);
178     release(&cons.lock);
```

10. fs.c(7)

```
kernel/fs.c

251     acquire(&itable.lock);
258         release(&itable.lock);
274     release(&itable.lock);
284     acquire(&itable.lock);
286     release(&itable.lock);
301     acquiresleep(&ip->lock);
326     releasesleep(&ip->lock);
339     acquire(&itable.lock);
345     // so this acquiresleep() won't block (or deadlock).
346     acquiresleep(&ip->lock);
347
348     release(&itable.lock);
349
355     releasesleep(&ip->lock);
356
357     acquire(&itable.lock);
358 }
361     release(&itable.lock);
```

11. uart.c(2 locks)

```
kernel/uart.c

89     acquire(&uart_tx_lock);
103    release(&uart_tx_lock);
185    }
186
187    // send buffered characters.
188    acquire(&uart_tx_lock);
189    uartstart();
190    release(&uart_tx_lock);
191 }
```

12. printf.c(1 lock)

```
kernel/printf.c

71     if(locking)
72         acquire(&pr.lock);
73
156     if(locking)
157         release(&pr.lock);
158
```

13. trap.c(1 lock)

```
kernel/trap.c

164  clockintr()
165  {
166      if(cpuid() == 0){
167          acquire(&tickslock);
168          ticks++;
169          wakeup(&ticks);
170          release(&tickslock);
```

14. virtio_disk.c (2 lock)

```
kernel/virtio_disk.c

220  acquire(&disk.vdisk_lock);

290
291  release(&disk.vdisk_lock);
292  }

296  {
297      acquire(&disk.vdisk_lock);
298
326  release(&disk.vdisk_lock);
```

You can view all these instances of `acquire` in the following search results:

[Lock occurrences in xv6 source]

<https://github.com/search?q=repo%3Amit-pdos%2Fxv6-riscv+acquire+release&type=code>

For semaphores, `sleep` and `wakeup` functions appear
proc.c(2 times baqi tou sleep and wakeup are only defined here)

```
kernel/proc.c

23 // helps ensure that wakeups of wait()ing
338 wakeup(initproc);
373 // Parent might be sleeping in wait().
374 wakeup(p->parent);
433 sleep(p, &wait_lock); //DOC: wait-sleep
533 // regular process (e.g., because it calls sleep), and thus cannot
545 // Atomically release lock and sleep on chan.
546 // Reacquires lock when awakened.
547 void
548 sleep(void *chan, struct spinlock *lk)
549 {
550     struct proc *p = myproc();
551
555     // guaranteed that we won't miss any wakeup
556     // (wakeup locks p->lock),
559     acquire(&p->lock); //DOC: sleeplock1
562     // Go to sleep.
564     p->state = SLEEPING;
576     // Wake up all processes sleeping on chan.
579     wakeup(void *chan)
586     if(p->state == SLEEPING && p->chan == chan) {
606     if(p->state == SLEEPING){
607         // Wake process from sleep().
676     [SLEEPING] "sleep ",
```

Sleeplock.c (1 semaphore found)

```
kernel/sleeplock.c

1  // Sleeping locks
10 #include "sleeplock.h"
13 init_sleeplock(struct sleeplock *lk, char *name)
15     initlock(&lk->lk, "sleep lock");
22 acquiresleep(struct sleeplock *lk)
26     sleep(lk, &lk->lk);
34 releasesleep(struct sleeplock *lk)
39     wakeup(lk);
41 }
42
43 int
44 holdingsleep(struct sleeplock *lk)
45 {
46     int r;
47
```


Pipe.c(4 times)

```
kernel/pipe.c

8  #include "sleeplock.h"

64  wakeup(&pi->nread);

67  wakeup(&pi->nwrite);

86  return -1;

87  }

88  if(pi->nwrite == pi->nread + PIPESIZE){ //DOC: pipewrite-full
89      wakeup(&pi->nread);
90      sleep(&pi->nwrite, &pi->lock);
91  } else {
92      char ch;

99  wakeup(&pi->nread);

118  sleep(&pi->nread, &pi->lock); //DOC: piperead-sleep

127  wakeup(&pi->nwrite); //DOC: piperead-wakeup

Show less
```

Log.c(2 times)

```
kernel/log.c

6  #include "sleeplock.h"

22  // sleep until the last outstanding end_op() commits.

132     sleep(&log, &log.lock);

135     sleep(&log, &log.lock);

162     wakeup(&log);

167     // call commit w/o holding locks, since not allowed
168     // to sleep with locks.
169     commit();

171     log.committing = 0;
172     wakeup(&log);
173     release(&log.lock);
```

Virtio_disk.c(2 times)

```
kernel/virtio_disk.c

14  #include "sleeplock.h"

181     wakeup(&disk.free[0]);

232     sleep(&disk.free[0], &disk.vdisk_lock);

284     while(b->disk == 1) {
285         sleep(b, &disk.vdisk_lock);
286     }

320     b->disk = 0;    // disk is done with buf
321     wakeup(b);
322
```

Console.c(1 time)

```
kernel/console.c

17  #include "sleeplock.h"

95      }
96      sleep(&cons.r, &cons.lock);
97      }

171      cons.w = cons.e;
172      wakeup(&cons.r);
173      }
```

Uart.c(1 time)

```
kernel/uart.c

97      // wait for uartstart() to open up space in the buffer.
98      sleep(&uart_tx_r, &uart_tx_lock);
99      }

153      // maybe uartputc() is waiting for space in the buffer.
154      wakeup(&uart_tx_r);
155
```

You can explore the occurrences of `sleep` and `wakeup` in the following link:

<https://github.com/search?q=repo%3Amit-pdos%2Frv6-riscv+sleep+wakeup&type=code>

4. How can we use locks and semaphores to ensure a process 2 running before a process 1?
Ans)

PROCESS 1

Acquire(&s→lock) released in sleep

sleep (s, &s→lock)

(PROCESS 1 CODE HERE AFTER IT WAKES UP)

Release(&s→lock)

PROCESS 2

Acquire (Acquire(&s→lock)

(PROCESS 2 CODE HERE)

wakeup(s)

Release (Acquire(&s→lock)

THE ABOVE CODE ENSURE PROCESS 2 RUNS BEFORE PROCESS 1...

Explanation:

Process 1 (Parent):

- It acquires a lock and then goes to sleep, meaning it must be woken up by another process to continue execution. The sleep function releases that respective lock, allowing Process 2 (the child) to acquire it and perform its tasks. After process 2 wakes up process 1 it continues its work and then releases the lock.

Process 2 (Child):

- After being scheduled, it acquires the lock, performs its work, and calls wakeup() to wake up Process 1 and release its lock. Once Process 1 is awake, it continues its execution and releases its lock and completes execution.
This coordination ensures that Process 2 completes its work before Process 1.